

Vivak Kumar

vk2853@gmail.com | +91-7025420873 | <https://www.linkedin.com/in/vivak-kumar-48b049ab/> | Bengaluru, Karnataka

SUMMARY

Staff-level full-stack engineer with 10+ years of experience building and scaling AI-driven SaaS platforms end-to-end — exactly the kind of technical depth EverAI needs to reach 100M+ users. Proven expertise in Python, RAG pipelines, vector databases, LangGraph, and distributed systems, combined with strong frontend (React, JavaScript) and backend (Ruby on Rails, Elasticsearch, AWS) skills to independently own complex, narrative-driven product features from concept to production. Experienced in AI-first workflows, real-time data pipelines, and mentoring cross-functional engineering teams aligned with fast-paced, high-growth product goals.

SKILLS

Python, LangGraph, RAG (Retrieval-Augmented Generation), Vector Databases, Elasticsearch, AWS, React, JavaScript (ES6), jQuery, Ruby on Rails, Ruby, Distributed Systems, RDBMS (MySQL / PostgreSQL), Redis, Sidekiq, EmberJS, ViewJS, Stimulus, HTML5, CSS3, HAML, SASS, RSpec, Git, GitHub, GitHub Actions, Docker, Terraform, CI/CD, AppSignal, New Relic, Agile/Scrum, Test-Driven Development (TDD), Clean Code, Scalable Architecture, Performance Tuning

EXPERIENCE

Senior Software Engineer — STG Labs / Onclusive India Mar 2022 – Present

- Architected and scaled an AI-powered SaaS platform for PR intelligence using Ruby on Rails, MySQL, and Elasticsearch, supporting real-time content ingestion, classification, and retrieval across millions of articles.
- Built and optimized AI-first data pipelines for content categorization and semantic search, reducing search latency by 15% and improving relevance of retrieval results — directly analogous to RAG and vector-search workflows.
- Engineered dynamic frontend dashboards and narrative-driven data visualizations using React and JavaScript, enabling users to explore AI-generated insights, trends, and media metrics in real time.
- Migrated Elasticsearch from v5 to v7, improving distributed search performance, scalability, and indexing throughput for large-scale content corpora.
- Integrated multi-modal media sources (YouTube, podcasts) into the AI content pipeline, expanding the platform's generative content coverage and enriching the user experience.
- Improved API payload optimization and client-side rendering performance for large datasets, ensuring smooth, cinematic-quality data experiences across the product.
- Designed and implemented CI/CD pipelines using GitHub Actions and Docker on AWS, enabling faster, zero-downtime deployments and improved developer velocity across distributed teams.
- Collaborated cross-functionally with product and customer success teams to triage complex escalations and shape long-term platform architecture aligned with business goals.

Tech Lead — L&T; Technology Services Jan 2020 – Mar 2022

- Led a team of 6+ engineers in delivering AI-adjacent, data-intensive product features end-to-end — from concept and HLD/LLD architecture through to production deployment.
- Independently owned and shipped user-facing reporting modules with complex API design, backend data aggregation, and frontend behaviour, cutting delivery delays by 20%.
- Drove technical roadmap alignment across squads, facilitating sprint planning, backlog grooming, and unblock sessions to maintain high-velocity, narrative-driven product delivery.
- Produced comprehensive technical documentation (HLD/LLD) for distributed system components, ensuring scalable architecture decisions were well understood across the organization.
- Mentored engineers across cross-functional squads, fostering a culture of clean code, test-driven development, and ownership — recognized with two Best Employee Awards.

SDE II — Reflektive Labs Pvt Ltd Aug 2019 – Jan 2022

- Enhanced AI-adjacent monitoring pipelines and Redis-backed semantic search functionality, achieving a 10% performance gain through targeted distributed systems optimizations.
- Mentored junior engineers through design reviews, pair programming, and onboarding, building a strong engineering culture centered on scalable, AI-ready architecture.

Consultant — Ciber (Philips account) Jul 2018 – Jul 2019

- Migrated Lucene-based search infrastructure to PostgreSQL full-text search, increasing retrieval precision and laying groundwork for more advanced AI-first search pipelines.
- Built a custom SQL query layer for dynamic, generative report generation, enabling flexible, data-driven content output at scale.