

Vivak Kumar

vk2853@gmail.com | +91-7025420873 | <https://www.linkedin.com/in/vivak-kumar-48b049ab/> | Bengaluru, Karnataka

SUMMARY

Staff-level full-stack engineer with 10+ years of experience building and scaling end-to-end SaaS platforms, bringing strong hands-on expertise in Python, React, Ruby on Rails, Elasticsearch, and distributed systems — technologies directly aligned with EverAI's fast-paced, AI-driven product environment. Proven track record of independently managing production pipelines, owning features from inception to deployment, and collaborating cross-functionally to deliver high-impact digital experiences at scale. Eager to contribute to EverAI's mission of reaching 500M users by building robust, scalable backend and frontend systems that power the future of AI companionship.

SKILLS

Python, React, JavaScript (ES6), Ruby on Rails, Ruby, Elasticsearch, AWS, RAG, LangGraph, Vector DB, Distributed Systems, RDBMS, MySQL, PostgreSQL, jQuery, ViewJS, Redis, Sidekiq, EmberJS, Stimulus, HTML5, CSS3, HAML, SASS, RSpec, Git, GitHub, GitHub Actions, Docker, Terraform, CI/CD, AppSignal, New Relic, Agile/Scrum, Test-Driven Development (TDD), Clean Code, Scalable Architecture, Performance Tuning

EXPERIENCE

Senior Software Engineer — STG Labs / Onclusive India Mar 2022 – Present

- Developed and scaled a high-traffic SaaS platform for PR monitoring using Ruby on Rails, MySQL, and Elasticsearch, supporting real-time data pipelines and large-scale content ingestion — directly analogous to managing AI content delivery pipelines at scale.
- Built and enhanced interactive frontend dashboards and data visualizations using React and JavaScript, enabling users to analyze metrics, articles, and trends in real time across a fast-paced digital platform.
- Optimized API payloads and client-side rendering for large datasets, improving frontend performance and usability — mirroring the production pipeline efficiency required in high-growth AI product environments.
- Migrated Elasticsearch from v5 to v7, improving search performance and scalability to support growing user bases and content volumes.
- Implemented article categorization and topic-based indexing, reducing search latency by 15% — demonstrating hands-on experience with search and retrieval systems aligned with RAG and vector DB workflows.
- Integrated media sources including YouTube and podcast platforms, expanding the platform's multimedia content pipeline — relevant to AI video and audio content management.
- Designed and implemented CI/CD pipelines using GitHub Actions and Docker on AWS, enabling faster deployments with reduced downtime and improved developer productivity across distributed teams.
- Collaborated directly with product and customer success teams to triage complex escalations and shape platform architecture aligned with long-term business goals, demonstrating the cross-functional ownership valued at EverAI.

Tech Lead — L&T; Technology Services Jan 2020 – Mar 2022

- Led a team of 6+ engineers in a fast-paced, cross-functional environment, driving end-to-end feature delivery from idea inception through production — independently managing the full development pipeline much like EverAI's senior contributors.
- Owned delivery of user-facing reporting modules, collaborating on frontend UX, API design, and backend data aggregation to ensure high-performance digital experiences at scale.
- Proposed and implemented a Detailed Report feature that reduced processing delays by 20%, demonstrating a bias for measurable impact and continuous production optimization.
- Produced HLD/LLD technical documentation and architecture deliverables, ensuring scalable and maintainable distributed system design.
- Facilitated sprint planning, backlog grooming, and team unblock sessions, maintaining high-velocity delivery in an agile environment.
- Received two Best Employee Awards for delivery ownership, innovation, and cross-team leadership.

SDE II — Reflektive Labs Pvt Ltd Aug 2019 – Jan 2022

- Enhanced platform monitoring and Redis-backed search functionality, achieving a 10% performance gain — demonstrating hands-on expertise with distributed caching and real-time data retrieval systems.
- Actively mentored junior engineers through design reviews, pair programming, and onboarding, fostering a high-performance engineering culture in a collaborative team environment.

Consultant — Ciber (Philips account) Jul 2018 – Jul 2019

- Migrated Lucene-based search to PostgreSQL full-text search, increasing query precision and reliability for a large-scale enterprise data platform.
- Built a custom SQL query layer for dynamic report generation, enabling flexible, user-driven data exploration across complex RDBMS schemas.