

Vivak Kumar

vk2853@gmail.com | +91-7025420873 | <https://www.linkedin.com/in/vivak-kumar-48b049ab/> | Bengaluru, Karnataka

SUMMARY

Senior full-stack engineer and architect with 10+ years of experience designing, deploying, and scaling production AI-ready SaaS platforms end-to-end. Experienced technical decision-maker who has led small senior teams to ship ambitious systems 0→1—spanning distributed backends, RAG pipelines, vector search, LLM integrations, and enterprise-grade infrastructure—making me a strong fit for A.Team's invite-only network of elite AI engineers and architects. Proven partner to CTOs and product leaders, balancing deep technical ownership with business outcomes across LLMs, NLP, and ML infrastructure missions.

SKILLS

Python, LangGraph, RAG (Retrieval-Augmented Generation), Vector Databases, LLMs / NLP / AI Systems, Distributed Systems, Elasticsearch, AWS, Ruby on Rails, React, JavaScript (ES6), jQuery, Vue.js (VueJS), RDBMS (MySQL / PostgreSQL), Redis, Scalable Architecture, Docker, Terraform, CI/CD, GitHub Actions, Ruby, RSpec, HTML5, CSS3, HAML, SASS, Sidekiq, EmberJS, Stimulus, Git, GitHub, AppSignal, New Relic, Agile/Scrum, Test-Driven Development (TDD), Clean Code, Performance Tuning

EXPERIENCE

Senior Software Engineer — STG Labs / Onclusive India Mar 2022 – Present

- Architected and scaled a production SaaS platform for enterprise PR monitoring, owning end-to-end design from model-driven article categorization to distributed backend APIs using Ruby on Rails, MySQL, and Elasticsearch at scale.
- Designed and deployed RAG-adjacent search and categorization pipelines, integrating Elasticsearch vector search capabilities to surface relevant media content in real time across millions of articles—reducing search latency by 15%.
- Led the 0→1 build of real-time data visualization dashboards and analytics modules using React and JavaScript, enabling enterprise clients to analyze PR metrics, articles, and trends across large datasets.
- Drove a critical Elasticsearch migration from v5 to v7, re-architecting index mappings and query layers to improve search performance, scalability, and relevance for NLP-heavy workloads.
- Expanded the platform's AI data surface by architecting and shipping integrations for YouTube and podcast media sources, broadening the NLP corpus for downstream analysis.
- Owned CI/CD pipeline design using GitHub Actions and Docker on AWS, enabling zero-downtime production deployments and improving engineering velocity across a distributed team.
- Partnered directly with product leaders and customer success to scope, propose, and execute solutions for complex enterprise escalations, acting as the primary technical decision-maker for platform architecture.

Tech Lead — L&T; Technology Services Jan 2020 – Mar 2022

- Led a senior team of 6+ engineers in designing and delivering enterprise AI-adjacent reporting and data aggregation systems, driving cross-functional roadmap alignment between engineering, product, and business stakeholders.
- Acted as the technical decision-maker for architecture and API design across multiple squads, producing HLD/LLD documentation that governed platform scalability and long-term system health.
- Scoped, proposed, and executed a Detailed Report feature end-to-end, cutting data delivery delays by 20% through optimized backend aggregation pipelines and distributed query strategies.
- Mentored engineers across squads through design reviews, sprint planning, and backlog grooming—fostering a culture of ownership and engineering excellence.
- Received two Best Employee Awards for delivery ownership and technical innovation on mission-critical enterprise systems.

SDE II — Reflektive Labs Pvt Ltd Aug 2019 – Jan 2022

- Built and optimized Redis-backed distributed search and monitoring infrastructure, achieving a 10% performance improvement in production systems serving real-time user workloads.

- Drove engineering culture by mentoring junior engineers through pair programming, architecture reviews, and structured onboarding within a fast-paced, senior-led team.

Consultant — Ciber (Philips account) Jul 2018 – Jul 2019

- Architected a migration from Lucene-based full-text search to PostgreSQL full-text search, improving query precision and simplifying the distributed data retrieval stack for enterprise reporting.
- Designed and implemented a custom SQL query layer for dynamic, enterprise-grade report generation, enabling flexible data access patterns across large RDBMS datasets.